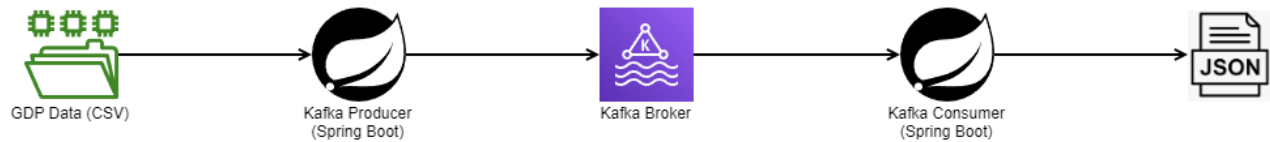


AWS POC 5 - Load Government Debt Total % GDP via Spring Boot & Kafka

Overview



Process Flow

1. Kafka Producer (Java) loads the CSV
2. Kafka Producer pushes the data record by record to the Kafka Broker
3. The Kafka Consumer (Java) listens and pulls the messages from the topic
4. The Kafka Consumer appends the record to the JSON file

Kafka Docker

docker-compose.yml

```
1  ---
2  version: '2'
3  services:
4
5    broker:
6      image: confluentinc/cp-kafka:7.6.0
7      hostname: broker
8      container_name: broker
9      ports:
10       - "9092:9092"
11       - "9101:9101"
12      environment:
13        KAFKA_NODE_ID: 1
14        KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: 'CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT'
15        KAFKA_ADVERTISED_LISTENERS: 'PLAINTEXT://broker:29092,PLAINTEXT_HOST://localhost:9092'
16        KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
17        KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0
18        KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1
19        KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1
20        KAFKA_JMX_PORT: 9101
21        KAFKA_JMX_HOSTNAME: localhost
22        KAFKA_PROCESS_ROLES: 'broker,controller'
23        KAFKA_CONTROLLER_QUORUM_VOTERS: '1@broker:29093'
24        KAFKA_LISTENERS: 'PLAINTEXT://broker:29092,CONTROLLER://broker:29093,PLAINTEXT_HOST://0.0.0.0:9092'
25        KAFKA_INTER_BROKER_LISTENER_NAME: 'PLAINTEXT'
26        KAFKA_CONTROLLER_LISTENER_NAMES: 'CONTROLLER'
27        KAFKA_LOG_DIRS: '/tmp/kraft-combined-logs'
28        # Replace CLUSTER_ID with a unique base64 UUID using "bin/kafka-storage.sh random-uuid"
29        # See https://docs.confluent.io/kafka/operations-tools/kafka-tools.html#kafka-storage-sh
30        CLUSTER_ID: 'MkU3OEVBNTcwNTJENDM2Qk'
```

```
31
32 schema-registry:
33   image: confluentinc/cp-schema-registry:7.6.0
34   hostname: schema-registry
35   container_name: schema-registry
36   depends_on:
37     - broker
38   ports:
39     - "8081:8081"
40   environment:
41     SCHEMA_REGISTRY_HOST_NAME: schema-registry
42     SCHEMA_REGISTRY_KAFKASTORE_BOOTSTRAP_SERVERS: 'broker:29092'
43     SCHEMA_REGISTRY_LISTENERS: http://0.0.0.0:8081
44
45 connect:
46   image: cnfldemos/cp-server-connect-datagen:0.6.4-7.6.0
47   hostname: connect
48   container_name: connect
49   depends_on:
50     - broker
51     - schema-registry
52   ports:
53     - "8083:8083"
54   environment:
55     CONNECT_BOOTSTRAP_SERVERS: 'broker:29092'
56     CONNECT_REST_ADVERTISED_HOST_NAME: connect
57     CONNECT_GROUP_ID: compose-connect-group
58     CONNECT_CONFIG_STORAGE_TOPIC: docker-connect-configs
59     CONNECT_CONFIG_STORAGE_REPLICATION_FACTOR: 1
60     CONNECT_OFFSET_FLUSH_INTERVAL_MS: 10000
61     CONNECT_OFFSET_STORAGE_TOPIC: docker-connect-offsets
62     CONNECT_OFFSET_STORAGE_REPLICATION_FACTOR: 1
63     CONNECT_STATUS_STORAGE_TOPIC: docker-connect-status
64     CONNECT_STATUS_STORAGE_REPLICATION_FACTOR: 1
65     CONNECT_KEY_CONVERTER: org.apache.kafka.connect.storage.StringConverter
66     CONNECT_VALUE_CONVERTER: io.confluent.connect.avro.AvroConverter
67     CONNECT_VALUE_CONVERTER_SCHEMA_REGISTRY_URL: http://schema-registry:8081
68     # CLASSPATH required due to CC-2422
69     CLASSPATH: /usr/share/java/monitoring-interceptors/monitoring-interceptors-7.6.0.jar
70     CONNECT_PRODUCER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringProducerInter"
71     CONNECT_CONSUMER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringConsumerInter"
72     CONNECT_PLUGIN_PATH: "/usr/share/java,/usr/share/confluent-hub-components"
73     CONNECT_LOG4J_LOGGERS: org.apache.zookeeper=ERROR,org.I0Itec.zkclient=ERROR,org.reflections=ERROR
74
75 control-center:
76   image: confluentinc/cp-enterprise-control-center:7.6.0
77   hostname: control-center
78   container_name: control-center
79   depends_on:
80     - broker
81     - schema-registry
82     - connect
83     - ksqldb-server
84   ports:
85     - "9021:9021"
86   environment:
87     CONTROL_CENTER_BOOTSTRAP_SERVERS: 'broker:29092'
88     CONTROL_CENTER_CONNECT_CONNECT-DEFAULT_CLUSTER: 'connect:8083'
```

```

89     CONTROL_CENTER_CONNECT_HEALTHCHECK_ENDPOINT: '/connectors'
90     CONTROL_CENTER_KSQL_KSQLDB1_URL: "http://ksqldb-server:8088"
91     CONTROL_CENTER_KSQL_KSQLDB1_ADVERTISED_URL: "http://localhost:8088"
92     CONTROL_CENTER_SCHEMA_REGISTRY_URL: "http://schema-registry:8081"
93     CONTROL_CENTER_REPLICATION_FACTOR: 1
94     CONTROL_CENTER_INTERNAL_TOPICS_PARTITIONS: 1
95     CONTROL_CENTER_MONITORING_INTERCEPTOR_TOPIC_PARTITIONS: 1
96     CONFLUENT_METRICS_TOPIC_REPLICATION: 1
97     PORT: 9021
98
99 ksqldb-server:
100     image: confluentinc/cp-ksqldb-server:7.6.0
101     hostname: ksqldb-server
102     container_name: ksqldb-server
103     depends_on:
104         - broker
105         - connect
106     ports:
107         - "8088:8088"
108     environment:
109         KSQL_CONFIG_DIR: "/etc/ksql"
110         KSQL_BOOTSTRAP_SERVERS: "broker:29092"
111         KSQL_HOST_NAME: ksqldb-server
112         KSQL_LISTENERS: "http://0.0.0.0:8088"
113         KSQL_CACHE_MAX_BYTES_BUFFERING: 0
114         KSQL_KSQL_SCHEMA_REGISTRY_URL: "http://schema-registry:8081"
115         KSQL_PRODUCER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringProducerInterce
116         KSQL_CONSUMER_INTERCEPTOR_CLASSES: "io.confluent.monitoring.clients.interceptor.MonitoringConsumerInterce
117         KSQL_KSQL_CONNECT_URL: "http://connect:8083"
118         KSQL_KSQL_LOGGING_PROCESSING_TOPIC_REPLICATION_FACTOR: 1
119         KSQL_KSQL_LOGGING_PROCESSING_TOPIC_AUTO_CREATE: 'true'
120         KSQL_KSQL_LOGGING_PROCESSING_STREAM_AUTO_CREATE: 'true'
121
122 ksqldb-cli:
123     image: confluentinc/cp-ksqldb-cli:7.6.0
124     container_name: ksqldb-cli
125     depends_on:
126         - broker
127         - connect
128         - ksqldb-server
129     entrypoint: /bin/sh
130     tty: true
131
132 ksql-datagen:
133     image: confluentinc/ksqldb-examples:7.6.0
134     hostname: ksql-datagen
135     container_name: ksql-datagen
136     depends_on:
137         - ksqldb-server
138         - broker
139         - schema-registry
140         - connect
141     command: "bash -c 'echo Waiting for Kafka to be ready... && \
142         cub kafka-ready -b broker:29092 1 40 && \
143         echo Waiting for Confluent Schema Registry to be ready... && \
144         cub sr-ready schema-registry 8081 40 && \
145         echo Waiting a few seconds for topic creation to finish... && \
146         sleep 11 && \

```

```

147         tail -f /dev/null'"
148     environment:
149         KSQL_CONFIG_DIR: "/etc/ksql"
150         STREAMS_BOOTSTRAP_SERVERS: broker:29092
151         STREAMS_SCHEMA_REGISTRY_HOST: schema-registry
152         STREAMS_SCHEMA_REGISTRY_PORT: 8081
153
154     rest-proxy:
155         image: confluentinc/cp-kafka-rest:7.6.0
156         depends_on:
157             - broker
158             - schema-registry
159         ports:
160             - 8082:8082
161         hostname: rest-proxy
162         container_name: rest-proxy
163         environment:
164             KAFKA_REST_HOST_NAME: rest-proxy
165             KAFKA_REST_BOOTSTRAP_SERVERS: 'broker:29092'
166             KAFKA_REST_LISTENERS: "http://0.0.0.0:8082"
167             KAFKA_REST_SCHEMA_REGISTRY_URL: 'http://schema-registry:8081'

```

Start the Kafka Docker Server

```
1 docker-compose up -d
```

Java Spring Boot Application

RestServiceApplication.java

```

1 @SpringBootApplication
2 public class RestServiceApplication {
3
4     public static void main(String[] args) {
5         SpringApplication.run(RestServiceApplication.class, args);
6     }
7
8 }

```

CountryGDPController.java

```

1 @RestController
2 @RequestMapping("/gdp")
3 public class CountryGDPController {
4     @Autowired
5     private CountryGDPMessagePublisher publisher;
6
7     @GetMapping("/is_running")
8     public GDPServiceStatus serviceStatus() {
9         return new GDPServiceStatus("GDP Service", "Running");
10    }
11
12    @GetMapping("/simulate_process_file")
13    public GDPDetailRecord simulateProcessFile() {
14        try {
15            var gdpDetailRecord = new GDPDetailRecord("Australia");
16            //publisher.sendGDPDetailRecordToTopic(gdpDetailRecord);
17            return gdpDetailRecord;
18        }
19    }
20 }


```

```

18     } catch (Exception ex) {
19         return new GDPDetailRecord("Failed to process");
20     }
21 }
22
23 }

```

 [marthinusswart/aws-proof-of-concepts](#)

 Created by marthinusswart • Updated 12 hours ago

Various Proof of Concept using the AWS Tech Stack

 Github

Kafka Docker Container

